A Minimal Process Model of Affective Metacognition: Gating (Threat/Safety × Defense), Contemplative Control, and Externalized Other-Model Updating (AI Outsourcing)

(anonymous draft)

Abstract

We propose a minimal process model of affective metacognition that bridges "brain–mind–body" dynamics with practical protocols for self-monitoring and AI-assisted cognition. The model combines (i) a gating mechanism driven by threat/safety appraisal (attachment) and defense processes, (ii) contemplative control that modulates prediction–error–salience via attention and interoception, and (iii) a label-as-hypothesis procedure: affect labels can down-regulate reactivity yet become harmful when reified. Finally, we treat intersubjectivity (shared reality) and AI outsourcing as the same class of externalized other-model updating, yielding testable trade-offs between cognitive relief and costs to memory, learning, and authorship/ownership.

1 Introduction

Figure 1 illustrates the proposed minimal process model and its operational protocol.

Affective metacognition—monitoring and steering one's emotional processes—often feels like a "between" space connecting brain, mind, and body. Existing literatures offer partial lenses: affect labeling reduces reactivity [9, 17], attachment organizes safety vs. threat regulation strategies [10, 14], mindfulness operationalizes attention and acceptance [2, 16], and predictive processing provides a unifying control language [6, 13]. Meanwhile, cognition is increasingly externalized to devices and LLMs; this resembles transactive memory and cognitive offloading [20, 12, 15], and may affect ownership of writing [8].

We unify these into a minimal model with four claims (C1–C4 in the accompanying state file). The goal is not a comprehensive theory, but a compact scaffold that (a) maps to measurable constructs, (b) yields practical self-protocols, and (c) supports design decisions for AI outsourcing without over-identifying with "tags."

2 Related Work (Very Brief)

Labeling and differentiation. Affect labeling can attenuate limbic responses [9] and is framed as implicit emotion regulation [17]. Yet verbal analysis can degrade preference quality in some contexts [21], motivating boundary conditions and careful label use. Emotion differentiation relates to regulation outcomes [1] and sits within broader process models of regulation [7].

Threat/safety and defenses. Attachment orientations are linked to hyperactivating vs. deactivating regulation strategies [10, 14]. Defense mechanisms admit hierarchical organization and measurement approaches [19, 18, 4]. Polyvagal theory provides a physiological lens on safety/threat states and behavioral repertoires [11].

Predictive processing and contemplative control. The free-energy principle frames perception/action as prediction error minimization [6]. Interoceptive inference connects bodily signals to felt emotion and selfhood [13]. Mindfulness is operationalized as attention regulation plus an accepting orientation [2], and has surveyed neural mechanisms [16].

Externalization, others, and AI. Shared reality describes motivation and mechanisms for aligning inner states [5]. Transactive memory shows distributed encoding in close relationships [20], and the Internet can behave like an external memory system [15]. Cognitive offloading formalizes when and why we outsource cognition [12]. The extended mind argues that external resources can be constitutive of cognition [3].

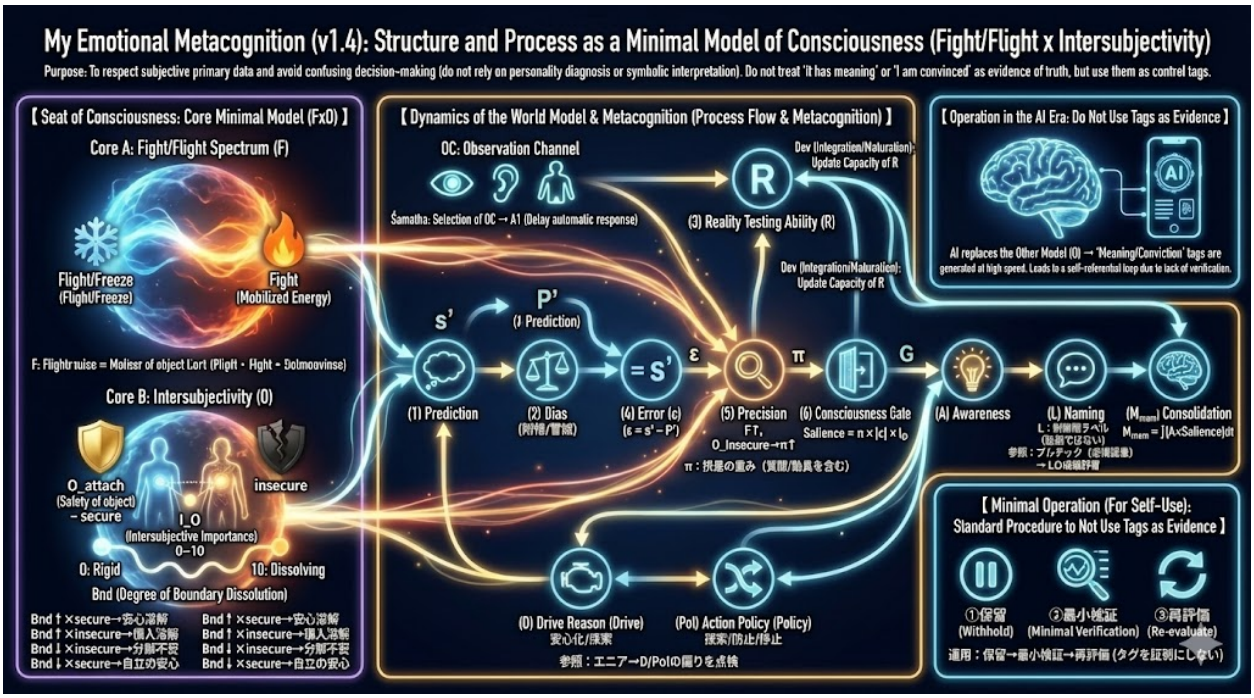

Figure 1: **Affective metacognition (v1.4):** a minimal process model integrating (i) threat/safety-driven gating (attachment $\times$ defenses), (ii) a prediction–error–salience loop modulated by contemplative observation (attention/interoception), (iii) a minimal "label-as-hypothesis" procedure (withhold $\rightarrow$ check $\rightarrow$ re-evaluate), and (iv) externalized other-model updating via social interaction and LLM outsourcing.

3 Model

3.1 Core variables and notation

We define a minimal set of latent state variables (scalars; extensions are possible):

- $T \in [0, 1]$: threat activation (higher = more threatened).
- $S = 1 - T$: perceived safety (higher = safer).
- D: dominant defense mode/level (e.g., mature $\rightarrow$ neurotic $\rightarrow$ immature).
- ε: prediction error (including interoceptive prediction error).
- σ: salience (attention weight) assigned to a channel/content.
- $G \in [0, 1]$: gate openness (access to exploration, labeling, reappraisal, flexible policy).

3.2 Gating: threat/safety $\times$ defense

We model a gating function:

$$G = f(S, D) \quad \text{with } \frac{\partial G}{\partial S} > 0, \tag{1}$$

and defense-dependent modulation (e.g., deactivation can keep G low by narrowing channels; hyperactivation can keep G noisy/high but unstable). Empirically, attachment insecurity correlates with deactivating vs. hyperactivating strategies [10]. Physiologically, safety/threat states reconfigure available behaviors [11]. Defenses provide a clinical language and hierarchy that can be measured [19, 4].

3.3 Contemplative control: prediction–error–salience loop

We assume salience allocation σ shapes how prediction error ε is amplified into felt urgency and subsequent policy selection. "Contemplative control" denotes deliberate regulation of σ (attention) and interoceptive sampling, consistent with operational definitions of mindfulness [2]. Predictive processing provides the control intuition: changing attention/precision alters the influence of prediction error [6, 13].

A minimal loop:

$$\text{(Observe)} \Rightarrow \varepsilon \tag{2}$$

$$\text{(Set precision / salience)} \Rightarrow \sigma \tag{3}$$

$$\text{(Update / choose policy)} \Rightarrow \pi \tag{4}$$

where mindfulness-like operations act primarily on σ (and secondarily on the generative model that produces ε).

3.4 Label-as-hypothesis: a minimal safe procedure

Affect labels can reduce reactivity [9, 17], but reification can freeze exploration or distort preferences [21]. We propose a minimal procedure that treats labels as hypotheses:

Minimal Label Procedure (MLP)
(1) *Withhold*: pause; do not finalize identity-level tags.
(2) *Name*: propose 1–2 candidate labels as hypotheses.
(3) *Check*: verify against body/interoception and context; seek alternatives.
(4) *Re-evaluate*: rename or drop; keep uncertainty explicit.

This aims to preserve the benefits of labeling while reducing fixation.

3.5 Externalized other-model updating: human and AI as one class

Let O denote an "other model" used for calibration of beliefs/feelings (shared reality) [5]. Interaction with another person or an LLM both instantiate an external update channel:

$$\theta_{t+1} = \theta_t + \eta \cdot U(\text{dialogue with } O), \tag{5}$$

where θ are internal models (self, others, world). This channel overlaps with cognitive offloading and transactive memory: distributing storage/processing can reduce load [12, 20], but can shift what is remembered ("where" over "what") [15]. Recent work suggests LLM interaction can influence perceived ownership of writing [8], making it essential to specify when outsourcing is beneficial vs. costly.

4 Applications

Self-use (minimal). Use MLP only when G is sufficiently open; when T is high, prioritize safety restoration (downshift) and basic observation rather than fine-grained labeling.

Clinical/developmental connection (optional layer). Map recurring patterns of low G to attachment and defense profiles, and track shifts in D or T across time using validated instruments [4].

LLM outsourcing. Treat prompts and chat history as externalized memory/other-model. Design safeguards: (i) keep "authorial core" notes internal, (ii) periodically restate goals in one's own words, (iii) run MLP before committing to identity-laden labels suggested by the model.

5 Predictions (Testable)

- **P1 (Boundary condition):** labeling reduces immediate reactivity, but repeated identity-level labels increase fixation and reduce exploration, especially under high threat (low S).

- **P2 (Gate dependence):** mindfulness/attention manipulations improve reappraisal mainly when G is moderately open; under strong threat, effects are mediated by safety restoration.
- **P3 (Outsourcing trade-off):** AI/offloading improves task performance under load but reduces later internal recall and may alter authorship/ownership judgments.

6 Limitations

This is a minimal scaffold. It does not diagnose, and it abstracts complex constructs (attachment, defenses, predictive processing) into low-dimensional variables. Empirical work should specify measurement models and boundary conditions.

7 Conclusion

Affective metacognition can be made tractable by a compact control model: gating (threat/safety $\times$ defense), contemplative control (prediction–error–salience), and safe labeling (hypothesis procedure), extended to modern cognition via externalized other-model updating (human/AI).

References

[1] Lisa Feldman Barrett, James J. Gross, Todd C. Christensen, and Michael Benvenuto. Knowing what you're feeling and knowing what to do about it: mapping the relation between emotion differentiation and emotion regulation. *Cognition and Emotion*, 15(6):713–724, 2001. doi: 10.1080/02699930143000239.

[2] Scott R. Bishop, Mark Lau, Shauna Shapiro, Linda Carlson, Nicole D. Anderson, James Carmody, Zindel V. Segal, Susan Abbey, Michael Speca, Drew Velting, and Gerald Devins. Mindfulness: A proposed operational definition. *Clinical Psychology: Science and Practice*, 11(3):230–241, 2004. doi: 10.1093/clipsy/bph077.

[3] Andy Clark and David Chalmers. The extended mind. *Analysis*, 58(1):7–19, 1998. doi: 10.1093/analys/58.1.7.

[4] Mariagrazia Di Giuseppe and J. Christopher Perry. The hierarchy of defense mechanisms: Assessing defensive functioning with the defense mechanisms rating scales q-sort. *Frontiers in Psychology*, 12:718440, 2021. doi: 10.3389/fpsyg.2021.718440.

[5] Gerald Echterhoff and E. Tory Higgins. Shared reality: Construct and mechanisms. *Current Opinion in Psychology*, 23:iv–vii, 2018. doi: 10.1016/j.copsyc.2018.09.003.

[6] Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010. doi: 10.1038/nrn2787.

[7] James J. Gross. The emerging field of emotion regulation: An integrative review. *Review of General Psychology*, 2(3):271–299, 1998. doi: 10.1037/1089-2680.2.3.271.

[8] Nataliya Kosmyna et al. Your brain on chatgpt: Accumulation of cognitive debt when using an ai assistant for essay writing task. *arXiv*, 2025. doi: 10.48550/arXiv.2506.08872.

[9] Matthew D. Lieberman, Naomi I. Eisenberger, Molly J. Crockett, Sabrina M. Tom, Jennifer H. Pfeifer, and Baldwin M. Way. Putting feelings into words: affect labeling disrupts amygdala activity in response to affective stimuli. *Psychological Science*, 18(5):421–428, 2007. doi: 10.1111/j.1467-9280.2007.01916.x.

[10] Mario Mikulincer, Phillip R. Shaver, and Galit Pereg. Attachment theory and affect regulation: The dynamics, development, and cognitive consequences of attachment-related strategies. *Motivation and Emotion*, 27(2):77–102, 2003. doi: 10.1023/A:1024515519160.

[11] Stephen W. Porges. The polyvagal perspective. *Biological Psychology*, 74(2):116–143, 2007. doi: 10.1016/j.biopsycho.2006.06.009.

[12] Evan F. Risko and Sam J. Gilbert. Cognitive offloading. *Trends in Cognitive Sciences*, 20(9):676–688, 2016. doi: 10.1016/j.tics.2016.07.002.

[13] Anil K. Seth. Interoceptive inference, emotion, and the embodied self. *Trends in Cognitive Sciences*, 17(11): 565–573, 2013. doi: 10.1016/j.tics.2013.09.007.

[14] Phillip R. Shaver and Mario Mikulincer. Adult attachment strategies and the regulation of emotion. In James J. Gross, editor, *Handbook of Emotion Regulation*, pages 446–465. Guilford Press, New York, 2007.

[15] Betsy Sparrow, Jenny Liu, and Daniel M. Wegner. Google effects on memory: cognitive consequences of having information at our fingertips. *Science*, 333(6043):776–778, 2011. doi: 10.1126/science.1207745.

[16] Yi-Yuan Tang, Britta K. Hölzel, and Michael I. Posner. The neuroscience of mindfulness meditation. *Nature Reviews Neuroscience*, 16(4):213–225, 2015. doi: 10.1038/nrn3916.

[17] Jared B. Torre and Matthew D. Lieberman. Putting feelings into words: Affect labeling as implicit emotion regulation. *Emotion Review*, 10(2):116–124, 2018. doi: 10.1177/1754073917742706.

[18] George E. Vaillant. *Ego Mechanisms of Defense: A Guide for Clinicians and Researchers*. American Psychiatric Pub, 1992. ISBN 0880484047.

[19] George E. Vaillant, Michael Bond, and Caroline O. Vaillant. An empirically validated hierarchy of defense mechanisms. *Archives of General Psychiatry*, 43(8):786–794, 1986. doi: 10.1001/archpsyc.1986.01800080072010.

[20] Daniel M. Wegner, Ralph Erber, and Paula Raymond. Transactive memory in close relationships. *Journal of Personality and Social Psychology*, 61(6):923–929, 1991. doi: 10.1037/0022-3514.61.6.923.

[21] Timothy D. Wilson and Jonathan W. Schooler. Thinking too much: introspection can reduce the quality of preferences and decisions. *Journal of Personality and Social Psychology*, 60(2):181–192, 1991. doi: 10.1037/0022-3514.60.2.181.